

# Applied Text Analytics & Natural Language Processing

with Dr. Mahdi Roozbahani  
& Wafa Louhichi

*Transformers*  
*Introduction (Attention) – Part 3*





# Attention

- Bahdanau et al. (2014) introduced the Attention mechanism to improve the performance of encoder-decoder models for Machine Translation
- Attention allows making predictions about the **output** by paying “attention” to some parts of the **input** sequence (we deal with **2 sequences**; an input and an output sequence)
- Attention mechanism fixed the following issues of sequence-to-sequence models:
  - Dealing with long-range dependencies (Dependency among words separated by multiple sentences) between words in a long sentence





# Self-Attention

- When creating the representation for each word in a sequence, Self-Attention tells the model to pay attention to which words in the same sequence .
- Example: in translating "I like exercising in the morning" , when we want to create a representation for the word "exercising":
  - We'd want to look at the subject "I" because languages usually have different verb forms for different subjects.
  - The other words don't matter that much in the translation of this word.
- Words can have different meanings when placed next to other words.



The cat drank the milk because **it** was hungry

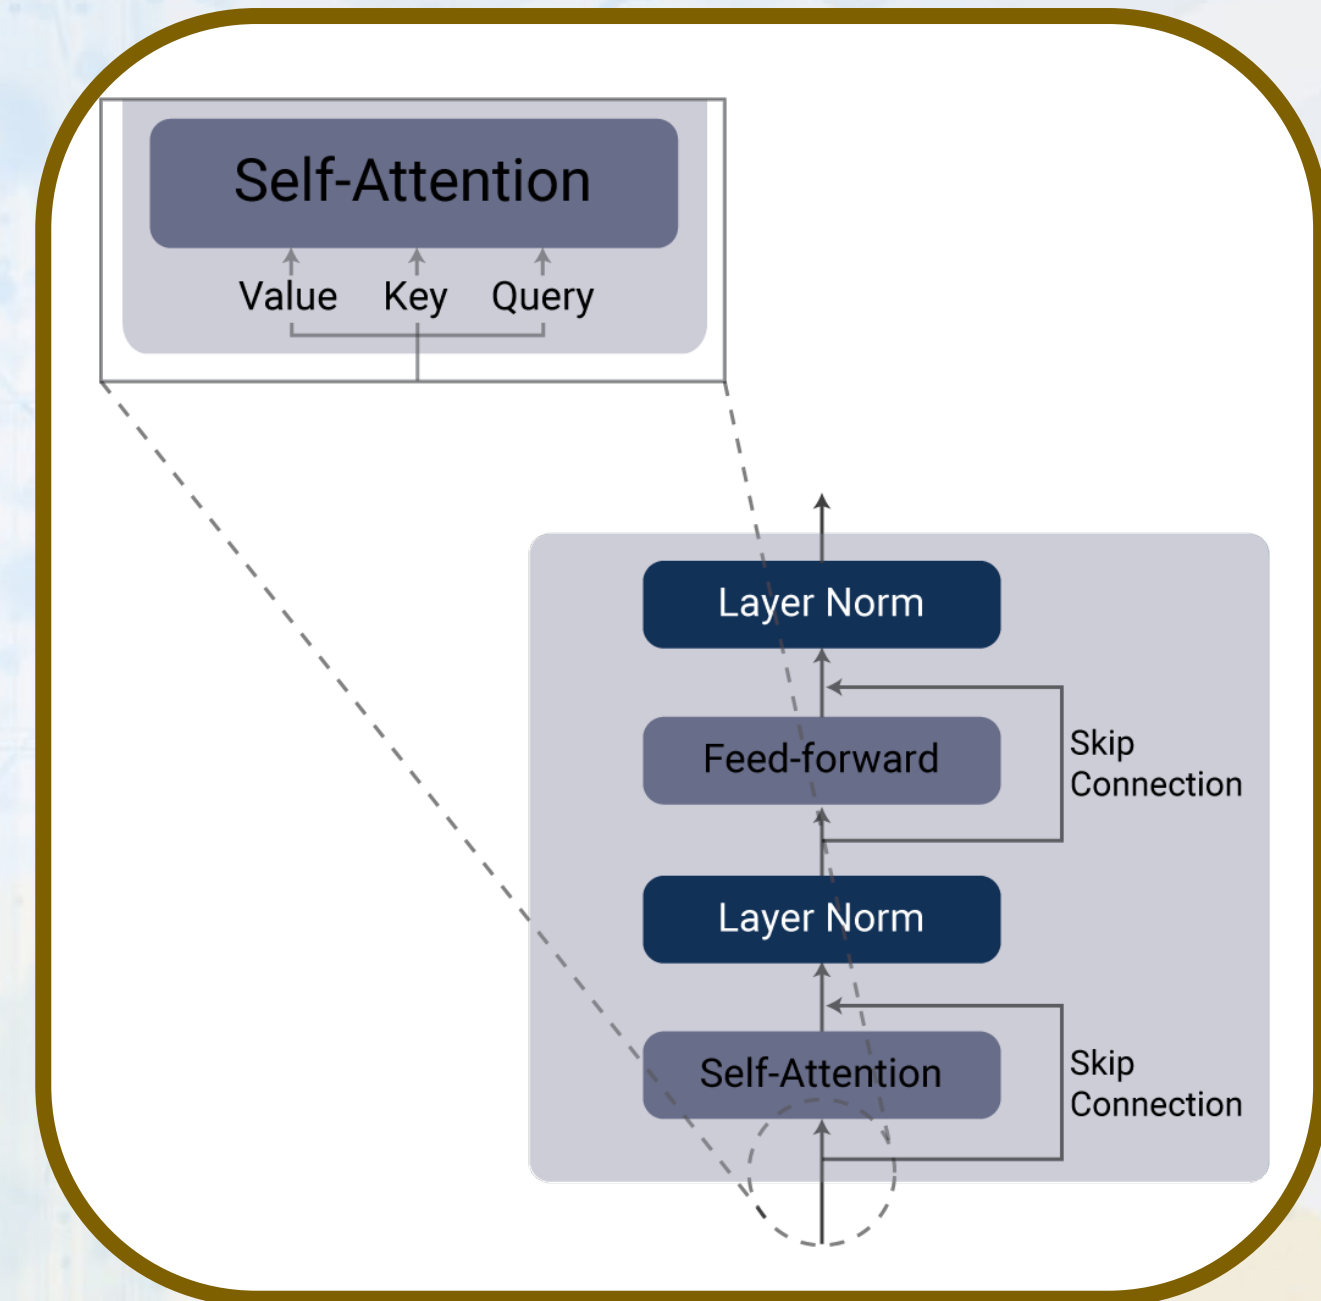


The cat drank the milk because **it** was sweet.

Self-attention extracts information about the meaning so that it can associate 'it' with the correct word.

# Self-Attention

- In the Self-Attention layer (in both encoder and decoder), the input is passed to three parameters: Key, Query and Value.
- The weighted values are computed by "a compatibility function of the query with the corresponding key".
- Key, Query and Value matrices are parameters whose initial weights are small, randomly selected numbers. These parameters change as the model is trained on the dataset.





# How to Calculate Attention for Each Input

- Each sequence (sentence) contains several words. The attention score must be calculated for each word.
  - First all the inputs need to be passed through embedding and positional encoding.
  - Key, Query and Value matrices need to be initialized randomly (their values will be optimized as the model is trained)
    - The size of all matrices are the same and depends on the number of words in a sequence and the length of embedding.
- Each vectorized word is multiplied by Key, Query and Value matrices to create key, query and value representation vector.

# How to Calculate Attention for Each Input

- Each sequence (sentence) contains several words. The attention score must be calculated for each word.
  - First all the inputs need to be passed through embedding and positional encoding.
  - Key, Query and Value matrices need to be initialized randomly (their values will be optimized as the model is trained)
    - The size of all matrices are the same and depends on the number of words in a sequence and the length of embedding.
- Each vectorized word is multiplied by Key, Query and Value matrices to create key, query and value representation vector.



# Calculate Attention Score

- This is where we can observe how self-attention is considering other words.
- This is the equation to compute each word's attention score where  $d_k$  is the embedding size, Q is the query matrix, K is the Key matrix and V is the Value matrix.

$$Attention(Q, K, V) = softmax\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$

# Calculate Attention Score

- This is where we can observe how self-attention is considering other words.
- This is the equation to compute each word's attention score where  $d_k$  is the embedding size, Q is the query matrix, K is the Key matrix and V is the Value matrix.

$$Attention(Q, K, V) = softmax\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$



# Summary

We learned about transformer models

- Encoder and decoder
- Why using transformer models
- Attention in transformers and why they are needed in such models

